Redes Software - Desarrollo SDN

Grupo 2
100451366 - Noel Andolz Aguado
100429717 - David Peño Gutiérrez

Copiar y pegar los ficheros de código fuente y poner fuente `Roboto Mono Normal - 10pt`

Hito 1:

- scenario.py → *vlanhost.py*
- simple_router.py → *L3Switch.py*
- Una captura de pantalla del terminal mininet con el comando `h1 ping -c4 h2`

*Los códigos de *vlanhost.py* y *L3Switch.py* se encuentran al final del hito

Primero ejecutamos mininet en un terminal.
Una vez vemos lo siguiente:

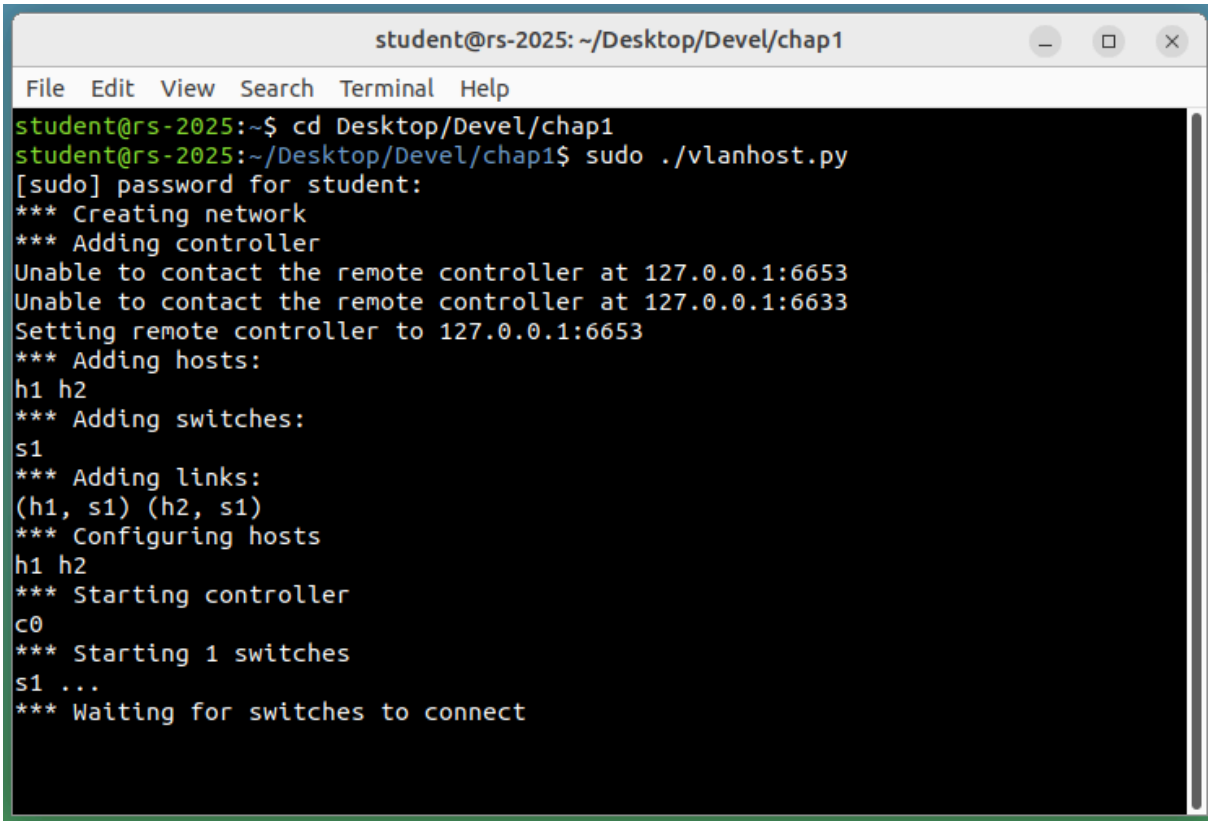

Abrimos Wireshark para poder ver la interacción entre el controlador y la topología de mininet.
Elegimos *any* para ver todo el tráfico. Y aplicamos el filtro "*openflow_v4*".
Vemos lo siguiente:

openflow_v4

No.	Time	Source	Destination	Protocol	Length	Info
105	49.329151469	127.0.0.1	127.0.0.1	OpenFl…	84	Type: OFPT_HELLO
107	49.331035013	127.0.0.1	127.0.0.1	OpenFl…	76	Type: OFPT_HELLO
109	49.331065475	127.0.0.1	127.0.0.1	OpenFl…	76	Type: OFPT_FEATURES_REQUEST
111	49.830202180	127.0.0.1	127.0.0.1	OpenFl…	100	Type: OFPT_FEATURES_REPLY
112	49.832461358	127.0.0.1	127.0.0.1	OpenFl…	84	Type: OFPT_MULTIPART_REQUEST, OFPMP_PORT_DESC
114	49.832912341	127.0.0.1	127.0.0.1	OpenFl…	276	Type: OFPT_MULTIPART_REPLY, OFPMP_PORT_DESC
115	49.833170324	127.0.0.1	127.0.0.1	OpenFl…	140	Type: OFPT_FLOW_MOD
116	49.833205542	127.0.0.1	127.0.0.1	OpenFl…	140	Type: OFPT_FLOW_MOD
117	49.833221584	127.0.0.1	127.0.0.1	OpenFl…	204	Type: OFPT_FLOW_MOD
118	49.833236191	127.0.0.1	127.0.0.1	OpenFl…	204	Type: OFPT_FLOW_MOD
119	49.833249589	127.0.0.1	127.0.0.1	OpenFl…	148	Type: OFPT_FLOW_MOD
123	54.837760088	127.0.0.1	127.0.0.1	OpenFl…	76	Type: OFPT_ECHO_REQUEST
124	54.838352313	127.0.0.1	127.0.0.1	OpenFl…	76	Type: OFPT_ECHO_REPLY
127	59.839763749	127.0.0.1	127.0.0.1	OpenFl…	76	Type: OFPT_ECHO_REQUEST
128	59.840588411	127.0.0.1	127.0.0.1	OpenFl…	76	Type: OFPT_ECHO_REPLY

Podemos ver que en los dos primeros mensajes capturados, se está iniciando la conexión.
Después, vemos tanto el "*features request*" como el "*features reply*".
Después, vemos una serie de mensajes en los que se envían las reglas:

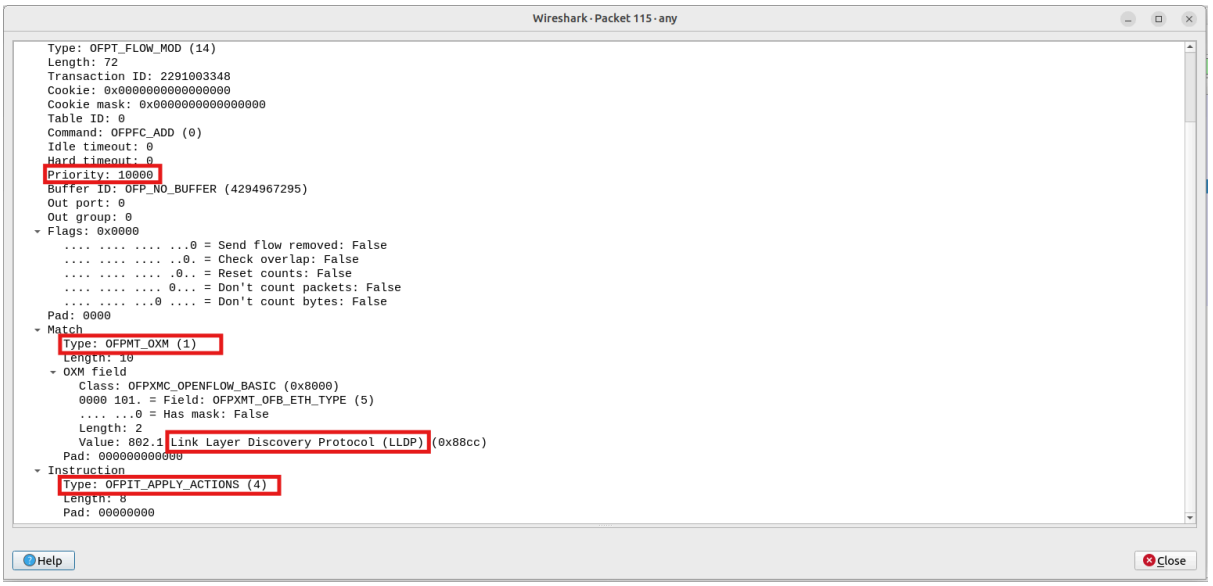

En esta primera, vemos las instrucciones para el tráfico LLDP, que tiene una prioridad de 10000 y que se descarta.

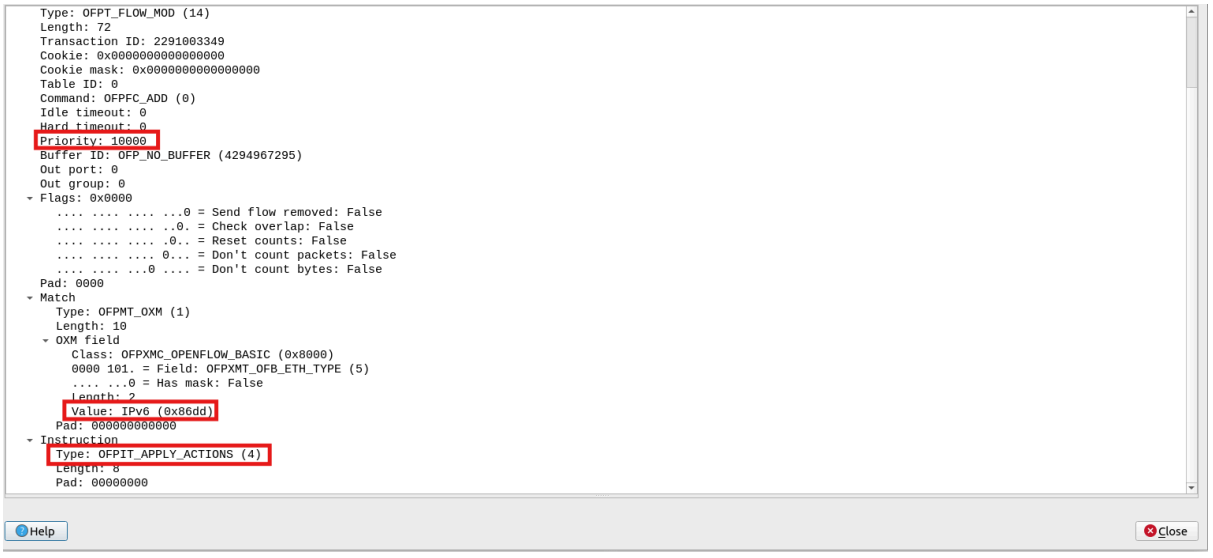

En este mensaje, podemos ver que se envían las instrucciones para el tráfico IPv6, que de nuevo tiene prioridad 10000 y también se descarta.

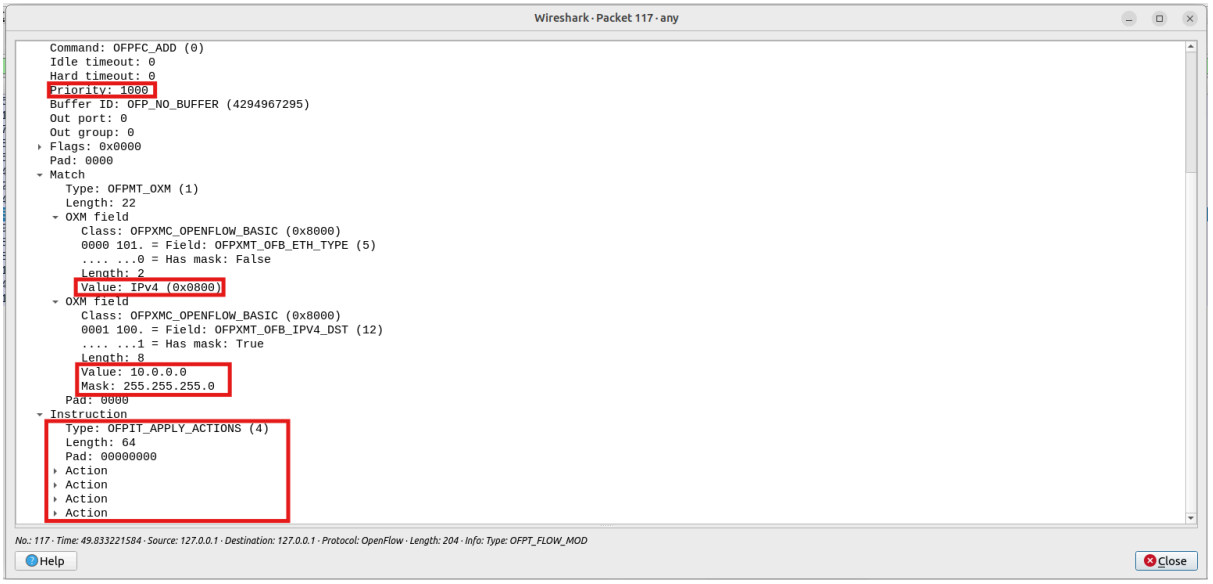

En este mensaje, vemos que se están enviando las reglas para el tráfico IPv4 cuyo prefijo es 10.0.0.0/24.

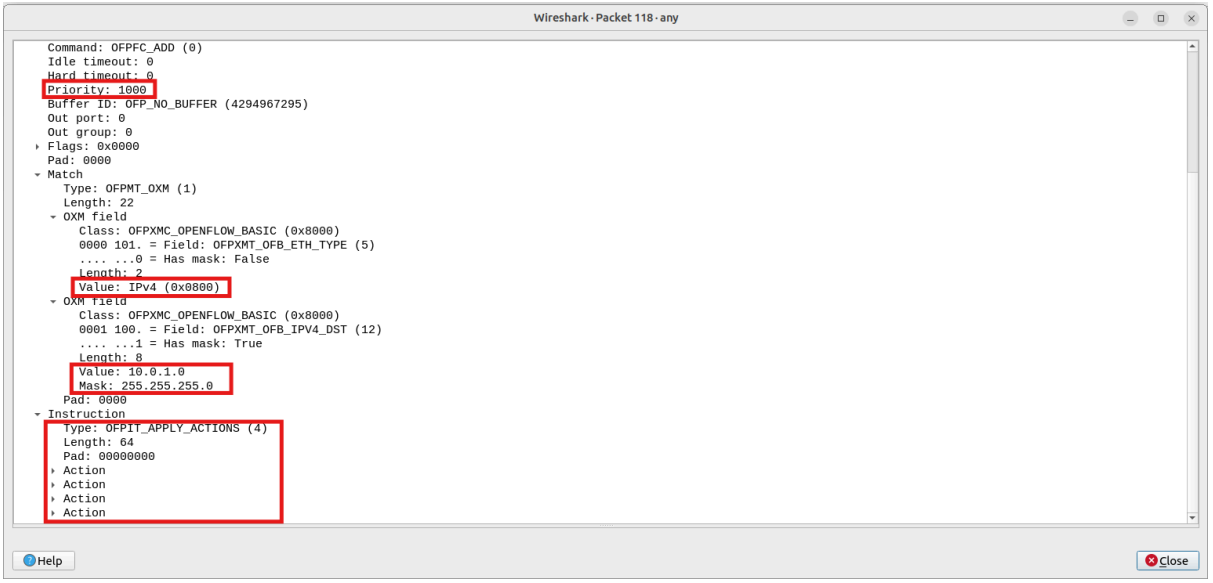

En este mensaje, vemos las reglas que se envían en el caso de tener tráfico IPv4 y cuyo prefijo sea 10.0.1.0/24.

Finalmente, tenemos el mensaje en el que se transmite la regla por defecto, indicando que se envía al controlador y con una prioridad de 0.

El siguiente tráfico que se puede ver son mensajes "*echo request*" y "*echo reply*" para asegurarse de que la conexión sigue establecida. Además, aprovechan para comprobar la latencia,...

Una vez vemos que el intercambio ha sido exitoso, podemos cerrar todo.

Volvemos a iniciar la topología, el controlador y Wireshark.

En Wireshark, en vez de capturar el tráfico en *any*, lo capturamos en *s1-eth1*.

Hacemos un ping desde h1 hasta h2:

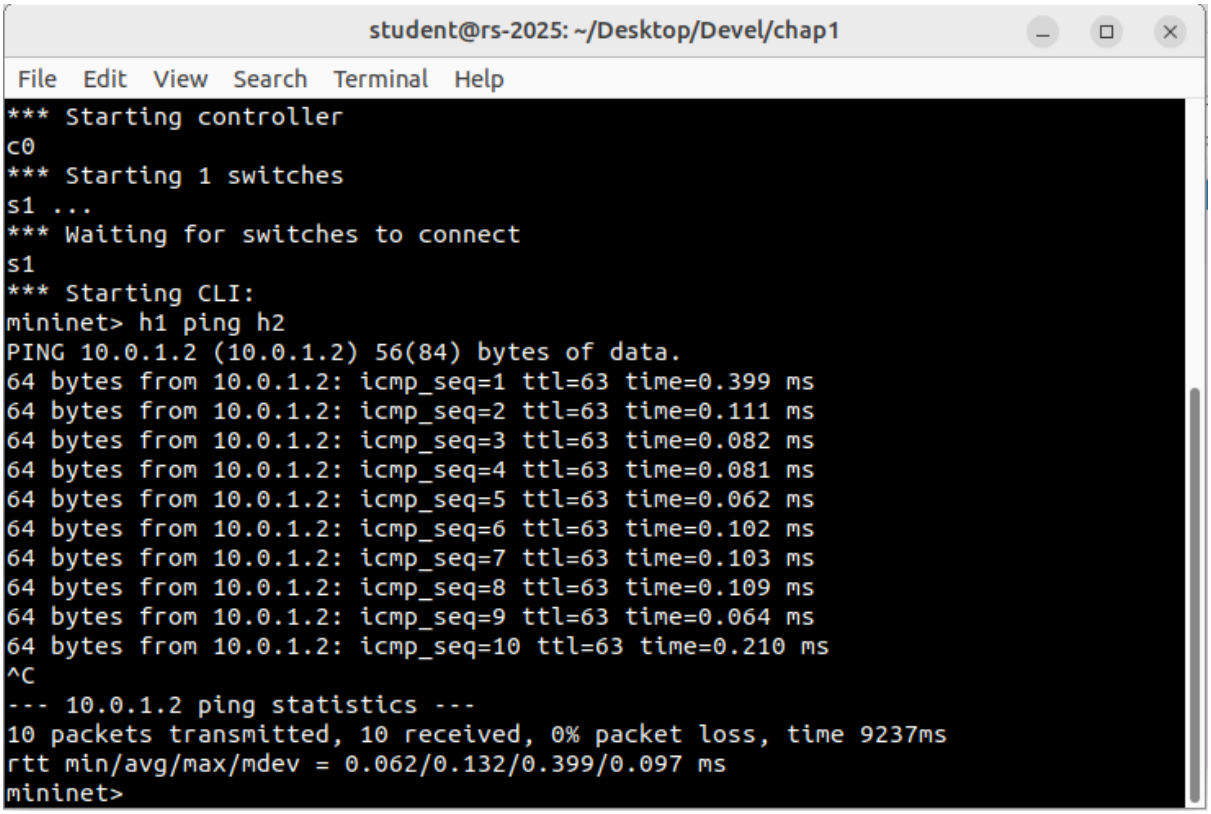

Y en Wireshark, vemos lo siguiente:

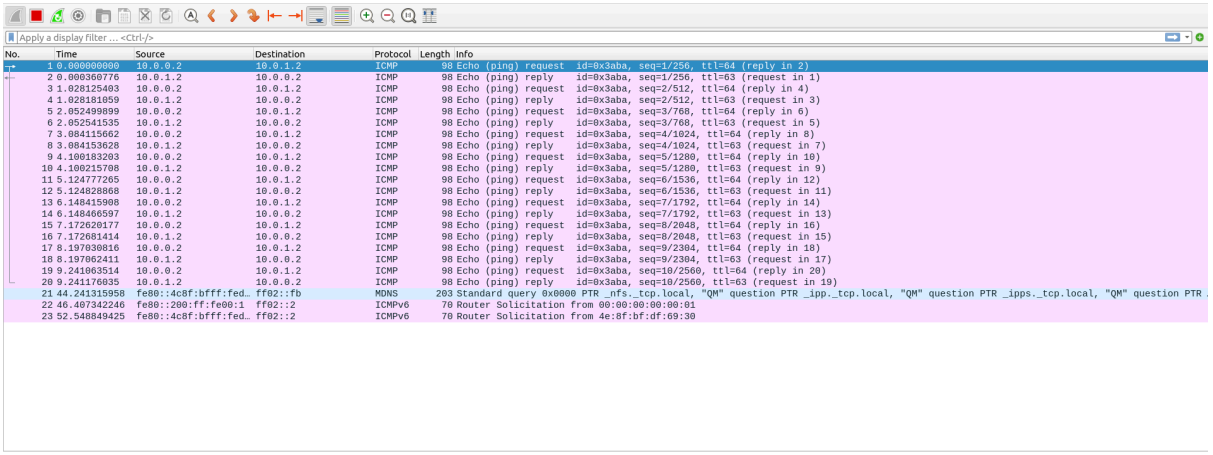

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=1/256, ttl=64 (reply in 2)
2	0.000360776	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=1/256, ttl=63 (request in 1)
3	1.028125403	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=2/512, ttl=64 (reply in 4)
4	1.028181059	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=2/512, ttl=63 (request in 3)
5	2.052499899	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=3/768, ttl=64 (reply in 6)
6	2.052541535	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=3/768, ttl=63 (request in 5)
7	3.084115662	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=4/1024, ttl=64 (reply in 8)
8	3.084153628	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=4/1024, ttl=63 (request in 7)
9	4.100183203	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=5/1280, ttl=64 (reply in 10)
10	4.100215708	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=5/1280, ttl=63 (request in 9)
11	5.124777265	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=6/1536, ttl=64 (reply in 12)
12	5.124828868	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=6/1536, ttl=63 (request in 11)
13	6.148415908	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=7/1792, ttl=64 (reply in 14)
14	6.148466597	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=7/1792, ttl=63 (request in 13)
15	7.172620177	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=8/2048, ttl=64 (reply in 16)
16	7.172681414	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=8/2048, ttl=63 (request in 15)
17	8.197030816	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=9/2304, ttl=64 (reply in 18)
18	8.197062411	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=9/2304, ttl=63 (request in 17)
19	9.241063514	10.0.0.2	10.0.1.2	ICMP	98	Echo (ping) request id=0x3aba, seq=10/2560, ttl=64 (reply in 20)
20	9.241176035	10.0.1.2	10.0.0.2	ICMP	98	Echo (ping) reply id=0x3aba, seq=10/2560, ttl=63 (request in 19)
21	44.241315958	fe80::4c8f:bfff:fed…	ff02::fb	MDNS	203	Standard query 0x0000 PTR _nfs._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question PTR _ipps._tcp.local, "QM" question PTR
22	46.407342246	fe80::200:ff:fe00:1	ff02::2	ICMPv6	70	Router Solicitation from 00:00:00:00:00:01
23	52.548849425	fe80::4c8f:bfff:fed…	ff02::2	ICMPv6	70	Router Solicitation from 4e:8f:bf:df:69:30

Nos metemos en los dos primeros mensajes para inspeccionarlos.
El primero:

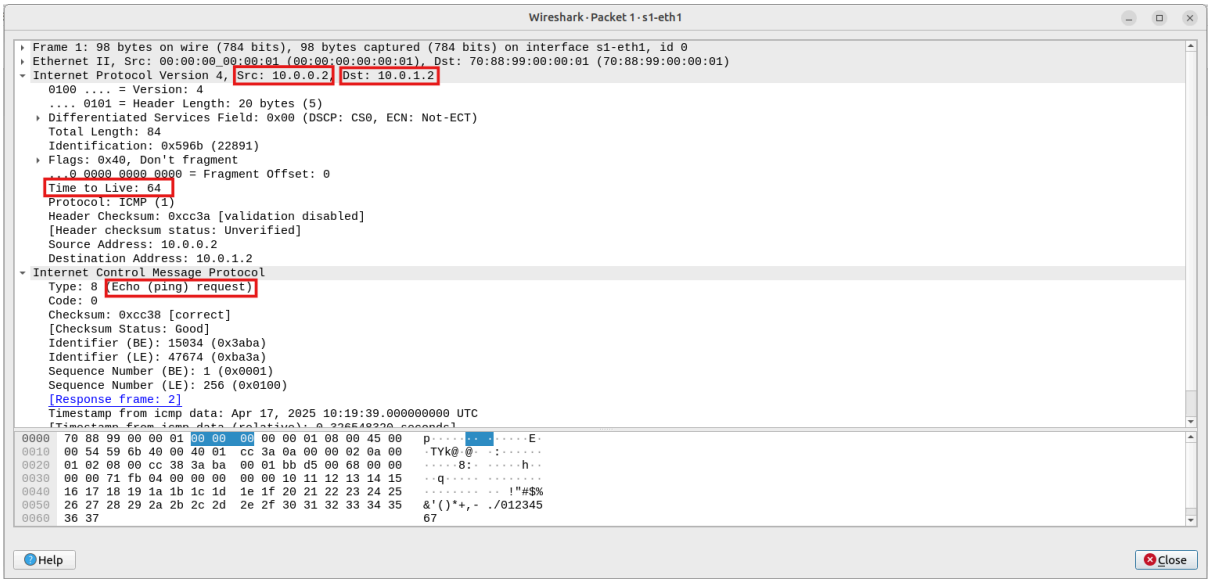

Vemos que es de tipo "*echo request*" y que lleva especificadas las direcciones IPv4 con origen en 10.0.0.2 y con destino 10.0.1.2. Además, podemos ver que el TTL es de 64.
Lo que esperamos ver en el siguiente paquete será un *echo reply* con dirección IPv4 de origen 10.0.1.2 y de destino 10.0.0.2. Además, el TTL debería ser 63.

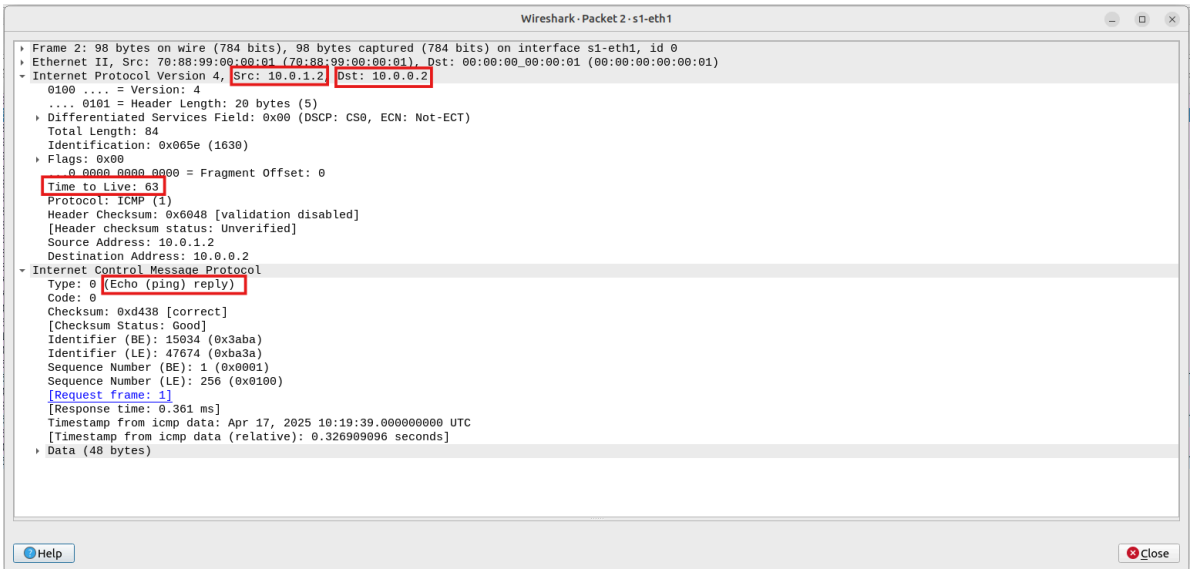

Efectivamente, lo que esperábamos ver es lo que hemos encontrado.

- **CÓDIGO vlanhost.py:**

```
#!/usr/bin/env python3
"""
vlanhost.py: Host subclass that uses a VLAN tag for the default interface.

Dependencies:
    This class depends on the "vlan" package
    $ sudo apt-get install vlan

Usage (example uses VLAN ID=1000):
    From the command line:
        sudo mn --custom vlanhost.py --host vlan,vlan=1000

    From a script (see exampleUsage function below):
        from functools import partial
        from vlanhost import VLANHost

        ....

        host = partial( VLANHost, vlan=1000 )
        net = Mininet( host=host, ... )

    Directly running this script:
        sudo python vlanhost.py 1000

"""

import sys
from mininet.node import Host
from mininet.topo import Topo
from mininet.util import quietRun
from mininet.log import error
from mininet.node import RemoteController, OVSSwitch
from mininet.net import Mininet
from mininet.cli import CLI
from mininet.topo import SingleSwitchTopo
from mininet.log import setLogLevel

from functools import partial

class ARPHost( Host ):
    "Host connected to ARP interface"
```

```
    # Alternatively, one could define
    # def config(self, vlan=None, **params)
    # and avoid the
    # vlan = params.pop('vlan' None)
    # pylint: disable=arguments-differ
    def config( self, **params ):
        """Configure ARPHost according to (optional) parameters:
           arp: ARP ID for default interface"""

        arp = params.pop('arp', None)
        assert arp is not None, 'ARPHost without arp in instantiation'
        mac, ip=arp
        r = super().config(**params)
        self.setARP(ip, mac)

        return r

class LineTopo(Topo):
    def build(self):
        switch = self.addSwitch('s1')
        h1=self.addHost('h1', cls=ARPHost,
arp=('70:88:99:00:00:01','10.0.0.1'),ip='10.0.0.2/24',
mac="00:00:00:00:00:01", defaultRoute=('via 10.0.0.1'))
        h2=self.addHost('h2', cls=ARPHost,
arp=('70:88:99:10:00:02','10.0.1.1'),ip='10.0.1.2/24',
mac="00:00:00:00:00:02", defaultRoute=('via 10.0.1.1'))
        self.addLink(h1, switch)
        self.addLink(h2, switch)

def exampleCustomTags():
    setLogLevel( 'info' )
    """Simple example that exercises LineTopo"""
    net = Mininet( topo=LineTopo(), waitConnected=True,
switch=partial(OVSSwitch, protocols="OpenFlow13"),
controller=partial(RemoteController, ip="127.0.0.1"))
    net.start()
    CLI(net)
    net.stop()


if __name__ == '__main__':
    exampleCustomTags()
```

- **CÓDIGO L3Switch.py:**

```
# Copyright (C) 2011 Nippon Telegraph and Telephone Corporation.
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
#    http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
# implied.
# See the License for the specific language governing permissions and
# limitations under the License.

from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
from ryu.lib.packet import ethernet
from ryu.lib.packet import ether_types
from ryu.ofproto import ether
from ryu.lib.packet import ipv4



class L3Switch(app_manager.RyuApp):
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]

    def __init__(self, *args, **kwargs):
        super(L3Switch, self).__init__(*args, **kwargs)
        self.mac_to_port = {}
        self.router_mac = {
            '10.0.0.1': '70:88:99:00:00:01',
            '10.0.1.1': '70:88:99:10:00:02',
        }

    def forward_actions(self, parser, ofproto, port, src_mac, dst_mac):
        """Acciones para reenviar un paquete con ajustes de cabecera L3."""
        return [
            parser.OFPActionDecNwTtl(),
            parser.OFPActionSetField(eth_src=src_mac),
```

```
            parser.OFPActionSetField(eth_dst=dst_mac),
            parser.OFPActionOutput(port),
        ]

    def drop_actions(self, parser, ofproto):
        """Acciones para descartar un paquete."""
        return []

    def send_to_controller_actions(self, parser, ofproto):
        """Acciones para enviar un paquete al controlador."""
        return [parser.OFPActionOutput(ofproto.OFPP_CONTROLLER)]




    @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
    def switch_features_handler(self, ev):
        datapath = ev.msg.datapath
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser

        # Quitar LLDP
        match = parser.OFPMatch(eth_type=ether.ETH_TYPE_LLDP)
        self.add_flow(datapath, 10000, match, self.drop_actions(parser,
ofproto))

        #Quitar trafico IPv6
        match = parser.OFPMatch(eth_type=ether.ETH_TYPE_IPV6)
        self.add_flow(datapath, 10000, match, self.drop_actions(parser,
ofproto))

        #IPv4
        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.0.0', '255.255.255.0')
        )
        actions = self.forward_actions(
            parser, ofproto, 1,
            '70:88:99:00:00:01',  # MAC origen
            '00:00:00:00:00:01'   # MAC destino = h1
        )

        self.add_flow(datapath, 1000, match, actions)

        # IPv4
```

```
        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.1.0', '255.255.255.0')
        )
        actions = self.forward_actions(
            parser, ofproto, 2,
            '70:88:99:10:00:02',  # MAC origen
            '00:00:00:00:00:02'   # MAC destino
        )
        self.add_flow(datapath, 1000, match, actions)

        # 3. Regla por defecto: enviar al controlador (prioridad 0)
        match = parser.OFPMatch()
        self.add_flow(datapath, 0, match,
self.send_to_controller_actions(parser, ofproto))




    def add_flow(self, datapath, priority, match, actions):
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser

        # construct flow_mod message and send it.
        inst = [parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
actions)]
        mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
match=match, instructions=inst)
        datapath.send_msg(mod)
```

Hito 2:

- Las modificaciones a simple_router.py
- Una captura de pantalla del terminal mininet con el comando `h2 ping -c4 10.0.1.1`

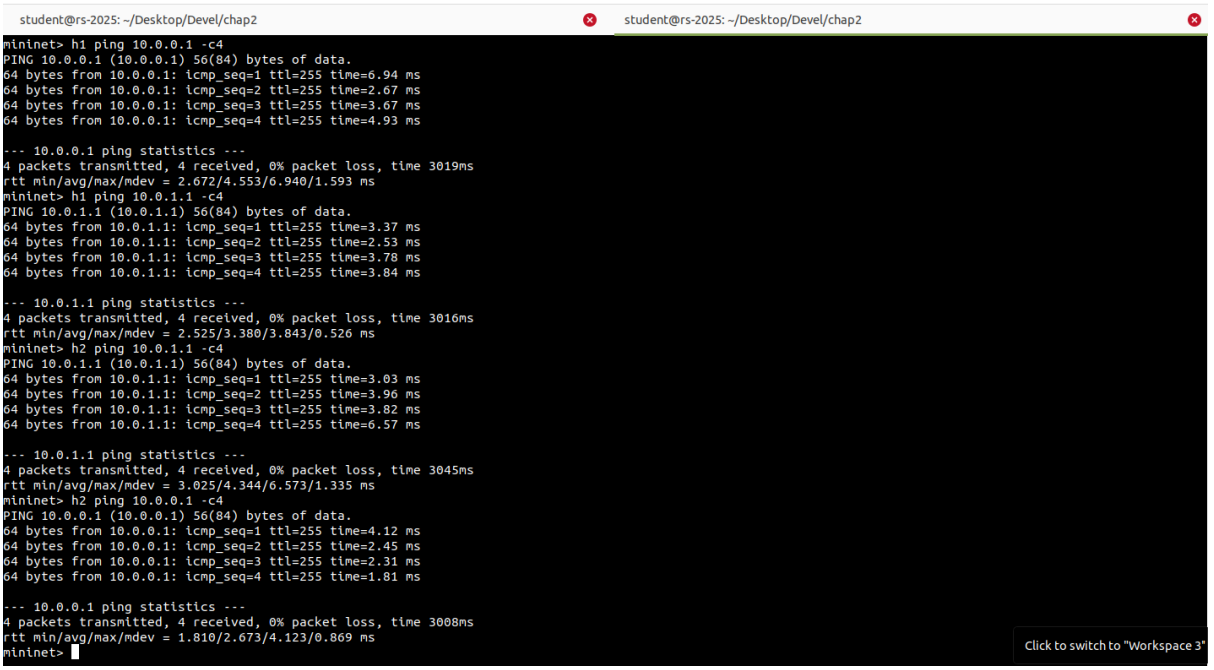

En esta captura de pantalla podemos ver los ping realizados desde el host 1 y host 2 hasta ambos puertos del encaminador. Efectivamente, comprobamos que los equipos finales pueden detectar ambos puertos del encaminador.

- **CÓDIGO L3Switch.py:**

```
from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
from ryu.lib.packet import ethernet
from ryu.lib.packet import ether_types
from ryu.ofproto import ether
from ryu.lib.packet import ipv4

from ryu.lib.packet import icmp
```

```
class L3Switch(app_manager.RyuApp):
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]

    def __init__(self, *args, **kwargs):
        super(L3Switch, self).__init__(*args, **kwargs)
        self.mac_to_port = {}
        self.router_mac = {
            '10.0.0.1': '70:88:99:00:00:01',
            '10.0.1.1': '70:88:99:10:00:02',
        }

    def forward_actions(self, parser, ofproto, port, src_mac,
dst_mac):
        """Acciones para reenviar un paquete con ajustes de cabecera
L3."""
        return [
            parser.OFPActionDecNwTtl(),
            parser.OFPActionSetField(eth_src=src_mac),
            parser.OFPActionSetField(eth_dst=dst_mac),
            parser.OFPActionOutput(port),
        ]

    def drop_actions(self, parser, ofproto):
        """Acciones para descartar un paquete."""
        return []

    def send_to_controller_actions(self, parser, ofproto):
        """Acciones para enviar un paquete al controlador."""
        return [parser.OFPActionOutput(ofproto.OFPP_CONTROLLER)]




    @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
    def switch_features_handler(self, ev):
        datapath = ev.msg.datapath
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser



        # Quitar LLDP
```

```
        match = parser.OFPMatch(eth_type=ether.ETH_TYPE_LLDP)
        self.add_flow(datapath, 10000, match,
self.drop_actions(parser, ofproto))

        #Quitar trafico IPv6
        match = parser.OFPMatch(eth_type=ether.ETH_TYPE_IPV6)
        self.add_flow(datapath, 10000, match,
self.drop_actions(parser, ofproto))


        match =parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.0.1', '255.255.255.255')
        )
        self.add_flow(datapath, 5000, match,
self.send_to_controller_actions(parser, ofproto))


        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.1.1', '255.255.255.255')
        )
        self.add_flow(datapath, 5000, match,
self.send_to_controller_actions(parser, ofproto))


        #IPv4
        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.0.0', '255.255.255.0')
        )
        actions = self.forward_actions(
            parser, ofproto, 1,
            '70:88:99:00:00:01',  # MAC origen
            '00:00:00:00:00:01'   # MAC destino = h1
        )

        self.add_flow(datapath, 1000, match, actions)

        # IPv4
        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.1.0', '255.255.255.0')
        )
        actions = self.forward_actions(
            parser, ofproto, 2,
```

```
            '70:88:99:10:00:02',  # MAC origen
            '00:00:00:00:00:02'   # MAC destino
        )
        self.add_flow(datapath, 1000, match, actions)

        # 3. Regla por defecto: enviar al controlador (prioridad 0)
        match = parser.OFPMatch()
        self.add_flow(datapath, 0, match,
self.send_to_controller_actions(parser, ofproto))

    @set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
    def _packet_in_handler(self, ev):
        msg = ev.msg
        datapath = msg.datapath
        port = msg.match['in_port']
        pkt = packet.Packet(data=msg.data)
        self.logger.info("packet-in %s" % (pkt,))
        pkt_ethernet = pkt.get_protocol(ethernet.ethernet)
        if not pkt_ethernet:
            return
        pkt_ipv4 = pkt.get_protocol(ipv4.ipv4)
        pkt_icmp = pkt.get_protocol(icmp.icmp)
        if pkt_icmp:
            self._handle_icmp(datapath, port, pkt_ethernet, pkt_ipv4,
pkt_icmp)
            return


    def _handle_icmp(self, datapath, port, pkt_ethernet, pkt_ipv4,
pkt_icmp):
        if pkt_icmp.type != icmp.ICMP_ECHO_REQUEST:
            return

        pkt = packet.Packet()

pkt.add_protocol(ethernet.ethernet(ethertype=pkt_ethernet.ethertype,
                                           dst=pkt_ethernet.src,

src=self.router_mac.get(pkt_ipv4.dst)))
        pkt.add_protocol(ipv4.ipv4(dst=pkt_ipv4.src,
                                   src=pkt_ipv4.dst,
                                   proto=pkt_ipv4.proto))
        pkt.add_protocol(icmp.icmp(type_=icmp.ICMP_ECHO_REPLY,
                                   code=icmp.ICMP_ECHO_REPLY_CODE,
                                   csum=0,
                                   data=pkt_icmp.data))
```

```
        #self.send_to_controller_actions(parser, ofproto)
        self._send_packet(datapath,port,pkt)

    def add_flow(self, datapath, priority, match, actions):
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser

        # construct flow_mod message and send it.
        inst =
[parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS, actions)]
        mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
match=match, instructions=inst)
        datapath.send_msg(mod)

    def _send_packet(self, datapath, port, pkt):
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser
        pkt.serialize()
        self.logger.info("packet-out %s" % (pkt,))
        data = pkt.data
        actions = [parser.OFPActionOutput(port=port)]
        out = parser.OFPPacketOut(datapath=datapath,
                                  buffer_id=ofproto.OFP_NO_BUFFER,
                                  in_port=ofproto.OFPP_CONTROLLER,
                                  actions=actions,
                                  data=data)

        datapath.send_msg(out)
```

Hito 3:

- Las modificaciones a scenario.py
- Las modificaciones a simple_router.py
- Una captura de pantalla de Wireshark con el tráfico para el comando `h1 ping -c 2 h2`
- Una captura de pantalla del terminal mininet para dicho comando

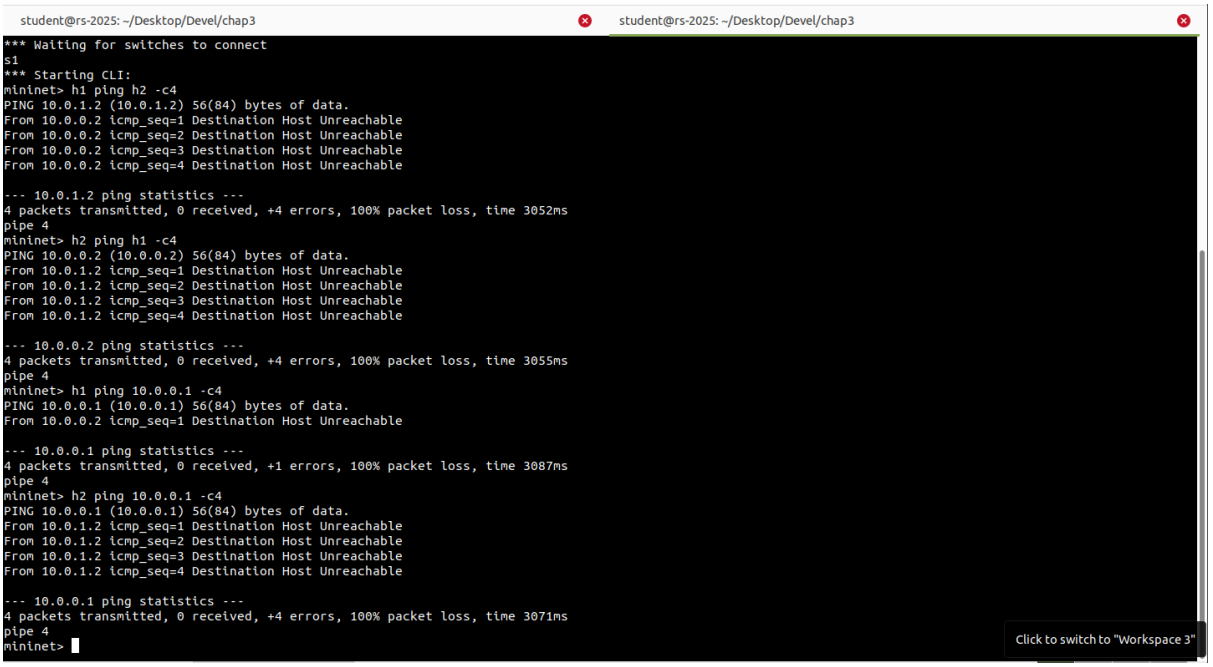

Al eliminar todo lo relacionado con la programación estática de la resolución de las direcciones MAC, es decir al eliminar todo lo relativo a ARP, y después, al ejecutar el controlador del capítulo anterior, vemos como no podemos hacer ping desde los equipos a cualquier otra dirección IP del escenario, tal y como se indica en la primera parte del hito.

Tras modificar el controlador (nuestro *L3Switch.py*)

```
student@rs-2025: ~/Desktop/Devel/chap3                    student@rs-2025: ~/Desktop/Devel/chap3
mininet> h1 ping h2 -c4
PING 10.0.1.2 (10.0.1.2) 56(84) bytes of data.
64 bytes from 10.0.1.2: icmp_seq=1 ttl=63 time=4.14 ms
64 bytes from 10.0.1.2: icmp_seq=2 ttl=63 time=0.074 ms
64 bytes from 10.0.1.2: icmp_seq=3 ttl=63 time=0.118 ms
64 bytes from 10.0.1.2: icmp_seq=4 ttl=63 time=0.109 ms

--- 10.0.1.2 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3019ms
rtt min/avg/max/mdev = 0.074/1.109/4.136/1.747 ms
mininet> h2 ping h1 -c4
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=63 time=0.227 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=63 time=0.108 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=63 time=0.084 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=63 time=0.058 ms

--- 10.0.0.2 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3072ms
rtt min/avg/max/mdev = 0.058/0.119/0.227/0.064 ms
mininet> h1 ping 10.0.0.1 -c4
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=255 time=4.10 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=255 time=2.77 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=255 time=3.32 ms
64 bytes from 10.0.0.1: icmp_seq=4 ttl=255 time=4.03 ms

--- 10.0.0.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3016ms
rtt min/avg/max/mdev = 2.770/3.556/4.100/0.546 ms
mininet> h2 ping 10.0.0.1 -c4
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=255 time=2.29 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=255 time=6.32 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=255 time=5.11 ms
64 bytes from 10.0.0.1: icmp_seq=4 ttl=255 time=2.61 ms

--- 10.0.0.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3011ms
rtt min/avg/max/mdev = 2.291/4.083/6.320/1.692 ms
mininet>
```

```
mininet> h1 ping 10.0.1.1 -c4
PING 10.0.1.1 (10.0.1.1) 56(84) bytes of data.
64 bytes from 10.0.1.1: icmp_seq=1 ttl=255 time=2.53 ms
64 bytes from 10.0.1.1: icmp_seq=2 ttl=255 time=3.13 ms
64 bytes from 10.0.1.1: icmp_seq=3 ttl=255 time=2.90 ms
64 bytes from 10.0.1.1: icmp_seq=4 ttl=255 time=3.60 ms

--- 10.0.1.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3006ms
rtt min/avg/max/mdev = 2.530/3.039/3.601/0.388 ms
mininet> h2 ping 10.0.1.1 -c4
PING 10.0.1.1 (10.0.1.1) 56(84) bytes of data.
64 bytes from 10.0.1.1: icmp_seq=1 ttl=255 time=2.99 ms
64 bytes from 10.0.1.1: icmp_seq=2 ttl=255 time=3.83 ms
64 bytes from 10.0.1.1: icmp_seq=3 ttl=255 time=3.59 ms
64 bytes from 10.0.1.1: icmp_seq=4 ttl=255 time=3.93 ms

--- 10.0.1.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3015ms
rtt min/avg/max/mdev = 2.992/3.585/3.927/0.364 ms
mininet>
```

Como podemos ver en las capturas, ahora sí es posible hacer ping a todos las direcciones.

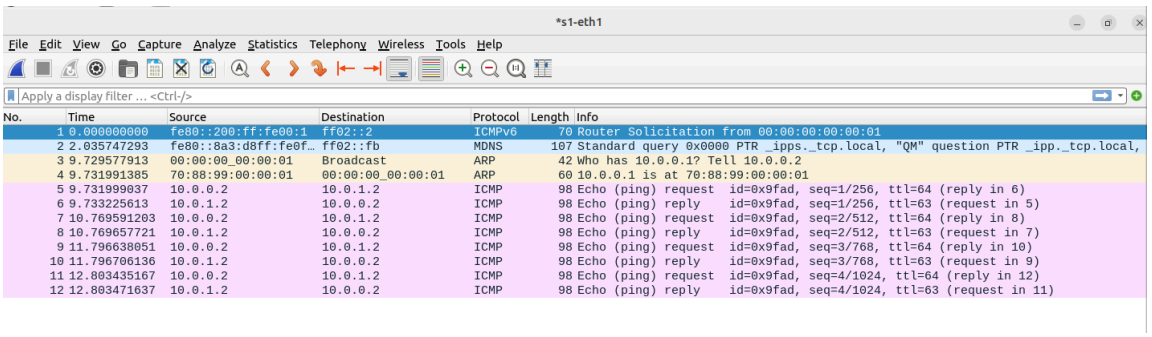

- **CÓDIGO vlanhost.py:**

```
#!/usr/bin/env python3

import sys
from mininet.node import Host
from mininet.topo import Topo
from mininet.util import quietRun
from mininet.log import error
from mininet.node import RemoteController, OVSSwitch
from mininet.net import Mininet
from mininet.cli import CLI
from mininet.topo import SingleSwitchTopo
from mininet.log import setLogLevel

from functools import partial



class LineTopo(Topo):
    def build(self):
        switch = self.addSwitch('s1')
        h1=self.addHost('h1', ip='10.0.0.2/24', mac="00:00:00:00:00:01",
defaultRoute=('via 10.0.0.1'))
        h2=self.addHost('h2', ip='10.0.1.2/24', mac="00:00:00:00:00:02",
defaultRoute=('via 10.0.1.1'))
        self.addLink(h1, switch)
        self.addLink(h2, switch)

def exampleCustomTags():
    setLogLevel( 'info' )
    net = Mininet( topo=LineTopo(), waitConnected=True,
switch=partial(OVSSwitch, protocols="OpenFlow13"),
controller=partial(RemoteController, ip="127.0.0.1"))
    net.start()
    CLI(net)
    net.stop()


if __name__ == '__main__':
    exampleCustomTags()
```

- **CÓDIGO L3Switch.py:**

```
from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
#from ryu.lib.packet import ethernet
from ryu.lib.packet import ether_types
from ryu.ofproto import ether
#from ryu.lib.packet import ipv4
#from ryu.lib.packet import icmp
from ryu.lib.packet import ethernet, ipv4, icmp, arp


class L3Switch(app_manager.RyuApp):
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]

    def __init__(self, *args, **kwargs):
        super(L3Switch, self).__init__(*args, **kwargs)
        self.mac_to_port = {}
        self.router_mac = {
            '10.0.0.1': '70:88:99:00:00:01',
            '10.0.1.1': '70:88:99:10:00:02',
        }

    def forward_actions(self, parser, ofproto, port, src_mac,
dst_mac):
        """Acciones para reenviar un paquete con ajustes de cabecera
L3."""
        return [
            parser.OFPActionDecNwTtl(),
            parser.OFPActionSetField(eth_src=src_mac),
            parser.OFPActionSetField(eth_dst=dst_mac),
            parser.OFPActionOutput(port),
        ]

    def drop_actions(self, parser, ofproto):
        """Acciones para descartar un paquete."""
        return []

    def send_to_controller_actions(self, parser, ofproto):
        """Acciones para enviar un paquete al controlador."""
        return [parser.OFPActionOutput(ofproto.OFPP_CONTROLLER)]
```

```
    @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
    def switch_features_handler(self, ev):
        datapath = ev.msg.datapath
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser



        # Quitar LLDP
        match = parser.OFPMatch(eth_type=ether.ETH_TYPE_LLDP)
        self.add_flow(datapath, 10000, match,
self.drop_actions(parser, ofproto))

        #Quitar trafico IPv6
        match = parser.OFPMatch(eth_type=ether.ETH_TYPE_IPV6)
        self.add_flow(datapath, 10000, match,
self.drop_actions(parser, ofproto))


        match =parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.0.1', '255.255.255.255')
        )
        self.add_flow(datapath, 5000, match,
self.send_to_controller_actions(parser, ofproto))


        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.1.1', '255.255.255.255')
        )
        self.add_flow(datapath, 5000, match,
self.send_to_controller_actions(parser, ofproto))


        #IPv4
        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.0.0', '255.255.255.0')
        )
        actions = self.forward_actions(
```

```
            parser, ofproto, 1,
            '70:88:99:00:00:01',  # MAC origen
            '00:00:00:00:00:01'   # MAC destino = h1
        )

        self.add_flow(datapath, 1000, match, actions)

        # IPv4
        match = parser.OFPMatch(
            eth_type=ether.ETH_TYPE_IP,
            ipv4_dst=('10.0.1.0', '255.255.255.0')
        )
        actions = self.forward_actions(
            parser, ofproto, 2,
            '70:88:99:10:00:02',  # MAC origen
            '00:00:00:00:00:02'   # MAC destino
        )
        self.add_flow(datapath, 1000, match, actions)

        # 3. Regla por defecto: enviar al controlador (prioridad 0)
        match = parser.OFPMatch()
        self.add_flow(datapath, 0, match,
self.send_to_controller_actions(parser, ofproto))

    @set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
    def _packet_in_handler(self, ev):
        msg = ev.msg
        datapath = msg.datapath
        port = msg.match['in_port']
        pkt = packet.Packet(data=msg.data)
        self.logger.info("packet-in %s" % (pkt,))
        pkt_ethernet = pkt.get_protocol(ethernet.ethernet)
        if not pkt_ethernet:
            return
        pkt_arp = pkt.get_protocol(arp.arp)
        if pkt_arp:
            self._handle_arp(datapath, port, pkt_ethernet, pkt_arp)
            return
        pkt_ipv4 = pkt.get_protocol(ipv4.ipv4)
        pkt_icmp = pkt.get_protocol(icmp.icmp)
        if pkt_icmp:
            self._handle_icmp(datapath, port, pkt_ethernet, pkt_ipv4,
pkt_icmp)
            return
```

```
    def _handle_icmp(self, datapath, port, pkt_ethernet, pkt_ipv4,
pkt_icmp):
        if pkt_icmp.type != icmp.ICMP_ECHO_REQUEST:
            return

        pkt = packet.Packet()

pkt.add_protocol(ethernet.ethernet(ethertype=pkt_ethernet.ethertype,
                                           dst=pkt_ethernet.src,

src=self.router_mac.get(pkt_ipv4.dst)))
        pkt.add_protocol(ipv4.ipv4(dst=pkt_ipv4.src,
                                   src=pkt_ipv4.dst,
                                   proto=pkt_ipv4.proto))
        pkt.add_protocol(icmp.icmp(type_=icmp.ICMP_ECHO_REPLY,
                                   code=icmp.ICMP_ECHO_REPLY_CODE,
                                   csum=0,
                                   data=pkt_icmp.data))
        #self.send_to_controller_actions(parser, ofproto)
        self._send_packet(datapath,port,pkt)

    def _handle_arp(self, datapath, port, pkt_ethernet, pkt_arp):
        if pkt_arp.opcode != arp.ARP_REQUEST:
            return
        pkt = packet.Packet()

pkt.add_protocol(ethernet.ethernet(ethertype=pkt_ethernet.ethertype,
                                           dst=pkt_ethernet.src,

src=self.router_mac.get(pkt_arp.dst_ip))) ######
        pkt.add_protocol(arp.arp(opcode=arp.ARP_REPLY,

src_mac=self.router_mac.get(pkt_arp.dst_ip),
                                 src_ip=pkt_arp.dst_ip,
                                 dst_mac=pkt_arp.src_mac,
                                 dst_ip=pkt_arp.src_ip))
        self._send_packet(datapath, port, pkt)



    def add_flow(self, datapath, priority, match, actions):
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser

        # construct flow_mod message and send it.
```

```
        inst =
[parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS, actions)]
        mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
match=match, instructions=inst)
        datapath.send_msg(mod)

    def _send_packet(self, datapath, port, pkt):
        ofproto = datapath.ofproto
        parser = datapath.ofproto_parser
        pkt.serialize()
        self.logger.info("packet-out %s" % (pkt,))
        data = pkt.data
        actions = [parser.OFPActionOutput(port=port)]
        out = parser.OFPPacketOut(datapath=datapath,
                                  buffer_id=ofproto.OFP_NO_BUFFER,
                                  in_port=ofproto.OFPP_CONTROLLER,
                                  actions=actions,
                                  data=data)

        datapath.send_msg(out)
```